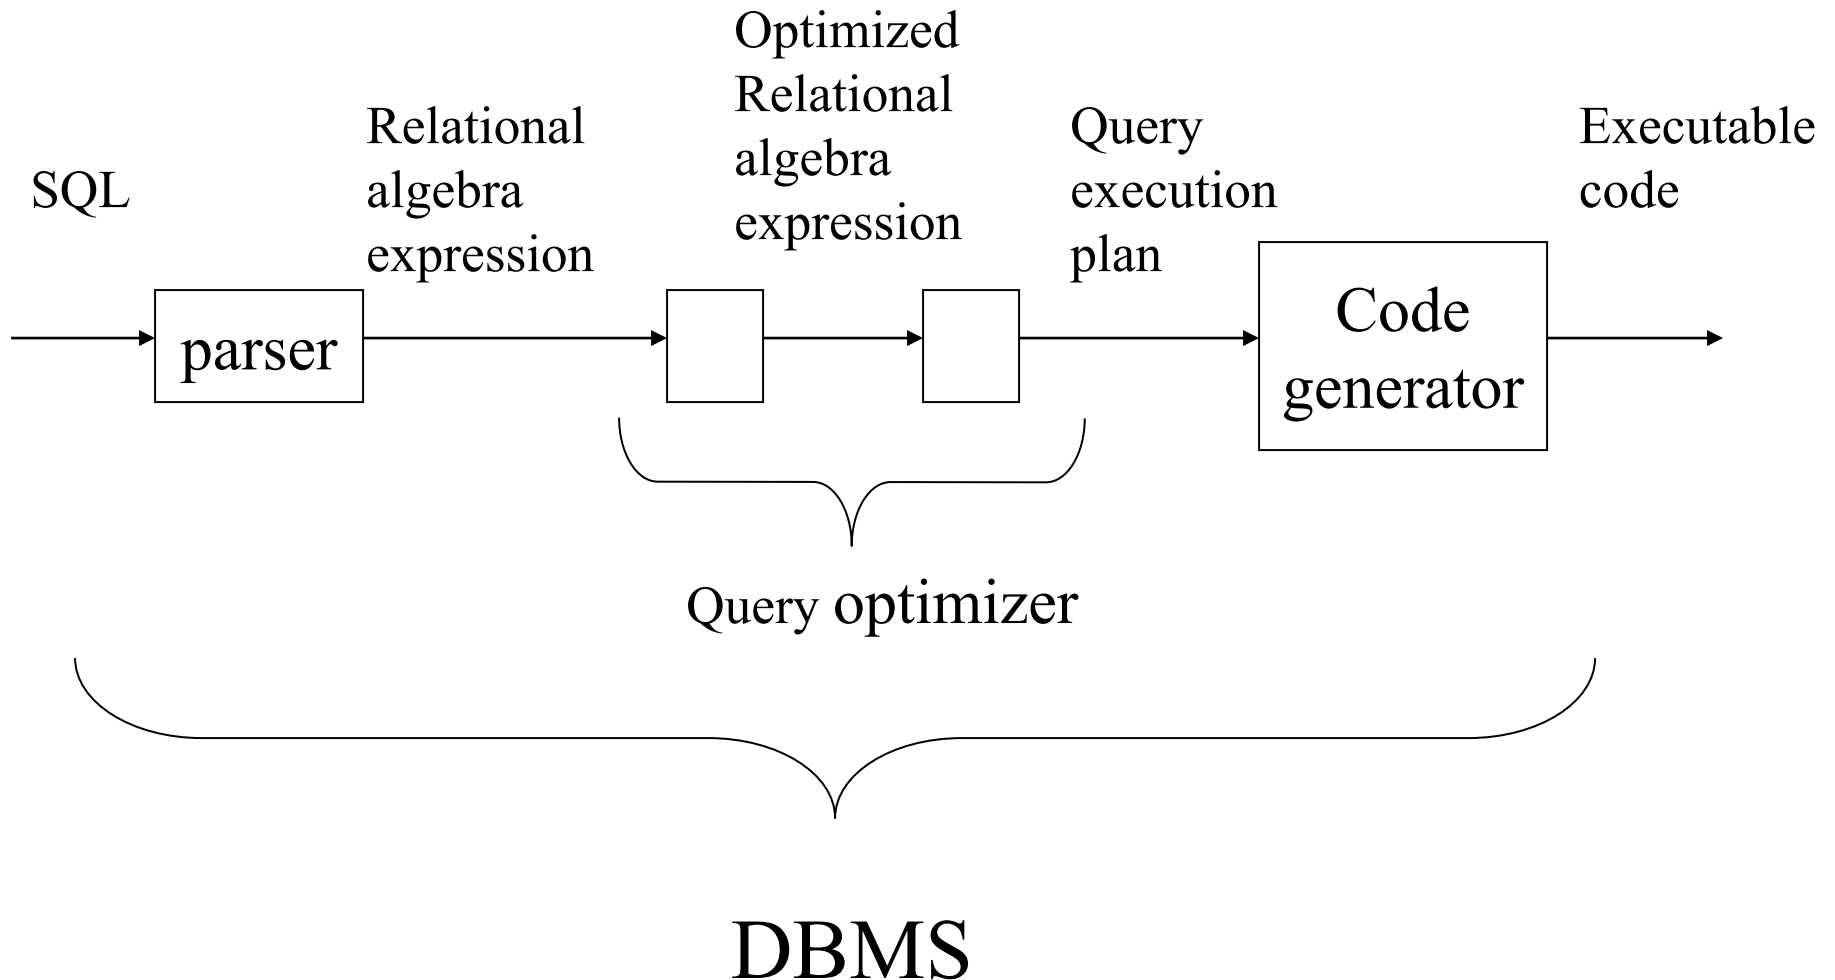


From Relational Algebra to SQL

W2013
CSCI 2141



Relational algebra, SQL, and the DBMS?



Basic Relational Algebra Operations (Unary): Selection, Projection

Selection: $\sigma_{\langle condition(s) \rangle} (\langle relation \rangle)$

- Picks tuples from the relation

Projection: $\Pi_{\langle attribute-list \rangle} (\langle relation \rangle)$

- Picks columns from the relation



Relational Algebra Operations (Set): Union, Set Difference



Union: ($\langle relation \rangle$) U ($\langle relation \rangle$)

- New relation contains all tuples from both relations, duplicate tuples eliminated.

Set Difference: $R - S$

- Produces a relation with tuples that are in R but NOT in S.



Relational Algebra Operations (Set): Cartesian Product, Intersect



Cartesian Product: $R \times S$

- Produces a relation that is concatenation of every tuple of R with every tuple of S
- The Above operations are the 5 fundamental operations of relational algebra.

Intersection: $R \cap S$

- All tuples that are in both R and S



Relational Algebra Operations (Join): Theta Join, Natural Join



Theta Join: $R \bowtie_F S = \sigma_F (R \times S)$

- Select all tuples from the Cartesian product of the two relations, matching condition F
- When F contains only equality “=”, it is called Equijoin

Natural Join: $R \bowtie S$

- Equijoin with common attributes eliminated

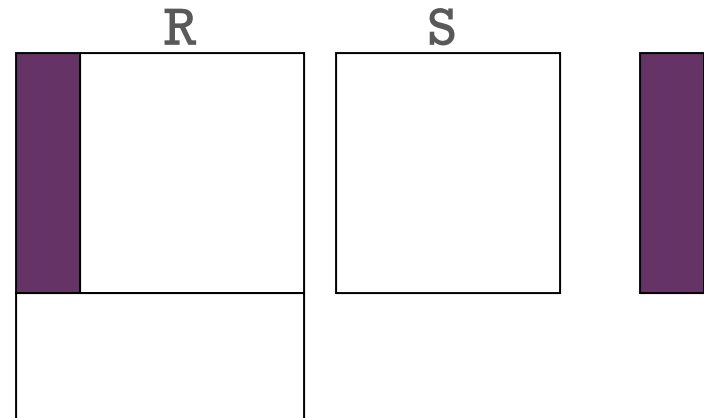


Relational Algebra Operations: Division



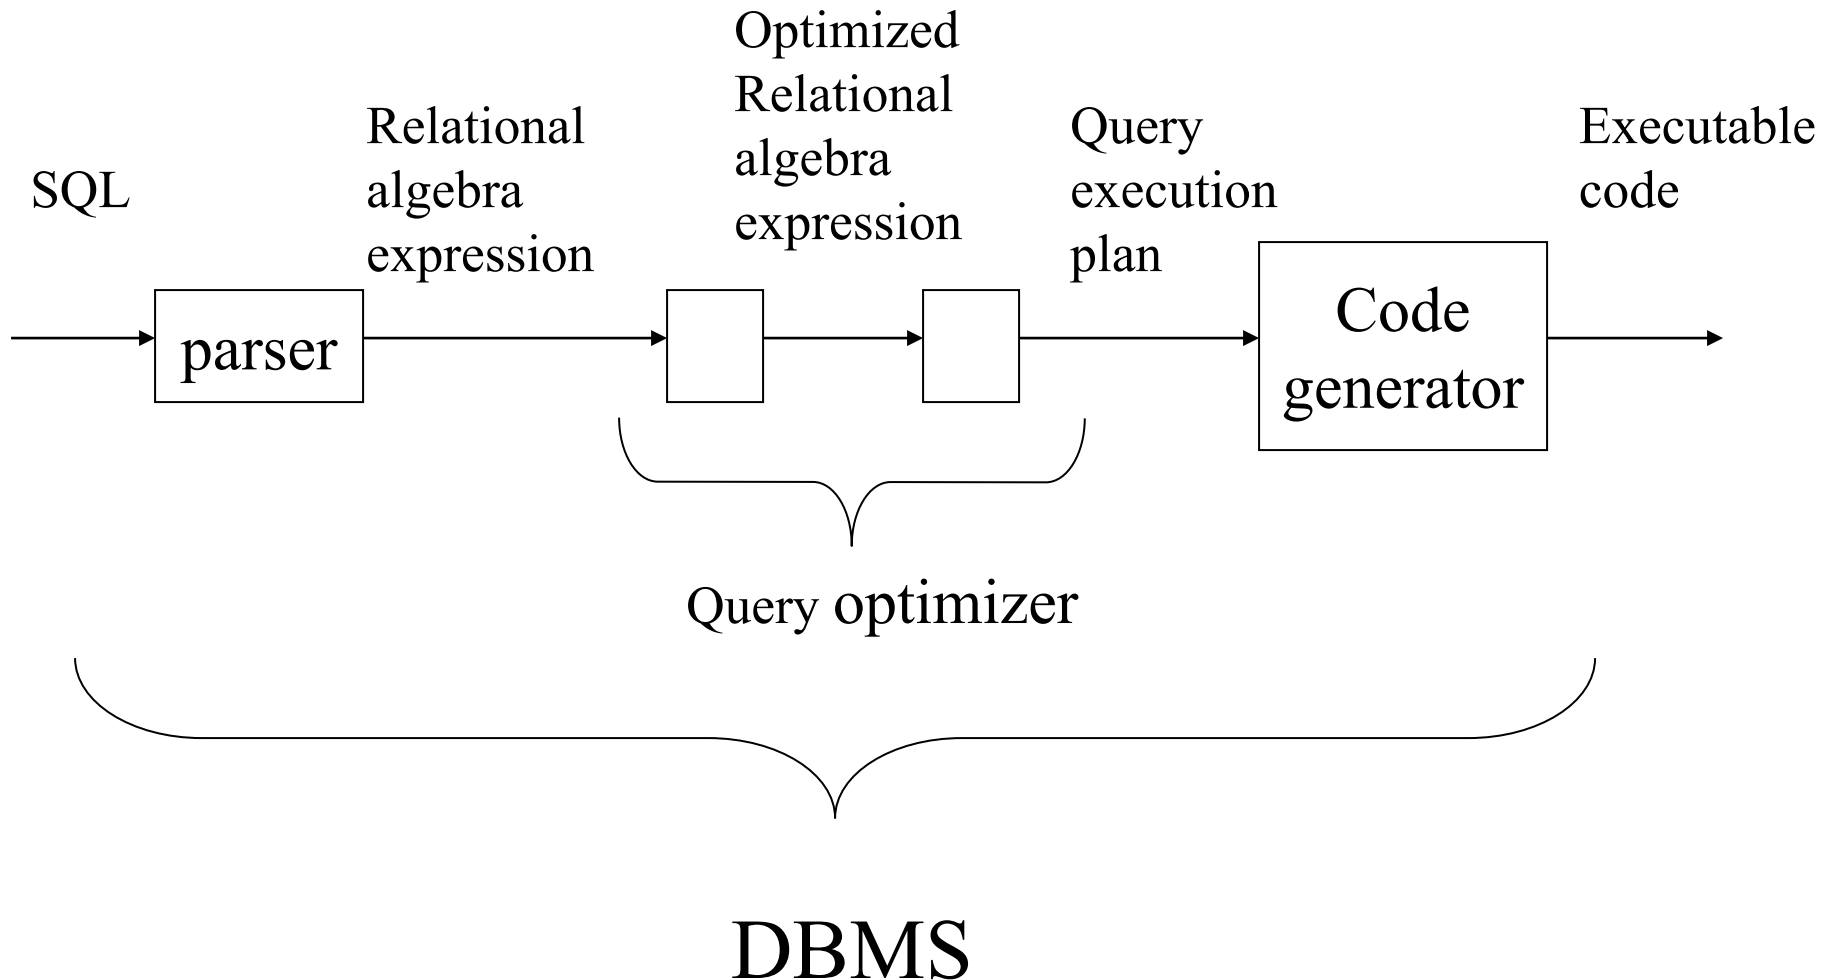
- Division: $R \div S$
- Produce a relation consist of the set of tuples from R that matches the combination of every tuple in S

$R \div S$





Relational algebra, SQL, and the DBMS?





SQL Introduction



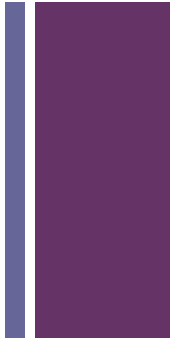
Standard language for querying and manipulating data

Structured Query Language

Many standards out there:

- ANSI SQL, SQL92 (a.k.a. SQL2), SQL99 (a.k.a. SQL3),
- Vendors support various subsets
- When starting to work with a specific database, make sure you are aware of the specific SQL form/syntax in use and the features supported
- MySQL:
 - <http://dev.mysql.com/doc/refman/5.0/en/what-is-mysql.html>
 - <http://dev.mysql.com/doc/refman/5.0/en/features.html>

+ SQL



- Data Definition Language (DDL)
 - Create/alter/delete tables and their attributes
 - <http://dev.mysql.com/doc/refman/5.0/en/tutorial.html>
 - 3.3.1 Creating and Selecting a Database
 - 3.3.2 Creating a Table
 - 3.3.3 Loading Data into a table
 - Friday
- Data Manipulation Language (DML)
 - Query one or more tables (today)
 - Insert/delete/modify tuples in tables (Monday)



Data Types in SQL



- Atomic types:
 - Characters: CHAR(20), VARCHAR(50), LONG VARCHAR
 - Numbers: INT, BIGINT, SMALLINT, FLOAT
 - Others: MONEY, DATETIME, ...
- MySQL
 - <http://dev.mysql.com/doc/refman/5.0/en/data-types.html>
 - Characters: CHAR(20), VARCHAR(50), MEDIUMTEXT
 - Numbers: INT, BIGINT, SMALLINT, FLOAT
 - Others: (couldn't find money), DATETIME
- Handy table of translations:
 - <http://dev.mysql.com/doc/refman/5.0/en/other-vendor-data-types.html>
- Every attribute must have an atomic type
 - Hence tables are flat



SQL Query



Basic form: (plus many many more bells and whistles)

```
SELECT <attributes>  
FROM   <one or more relations>  
WHERE  <conditions>
```



Relational algebra queries to SQL queries



- FROM clause produces Cartesian product (\times) of listed tables
- WHERE clause assigns rows to C in sequence and produces table containing only rows satisfying condition (sort of like σ)
- SELECT clause retains listed columns (Π)



Example Relational Algebra Translation to SQL

- `SELECT C.CrsName`
- `FROM Course C, Teaching T`
- `WHERE C.CrsCode=T.CrsCode AND T.Sem= 'W2013'`
- List CS courses taught in W2013
- Tuple variables (Course C, Teaching T) (aka aliases in MySQL) clarify meaning (Course AS C)
- Join condition “C.CrsCode=T.CrsCode”
 - eliminates those courses not being taught (junk)
- Selection condition “T.Sem= 'W2013' ”
 - eliminates irrelevant rows
- Equivalent (using natural join) to:

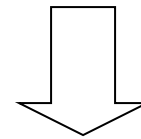
$\pi_{CrName}(Course \bowtie (\sigma_{Sem='W2013'}(Teaching)))$
 $\pi_{CrName}(\sigma_{Sem='W2013'}(Course \bowtie Teaching))$

+ Simple SQL Query

Product

PName	Price	Category	Manufacturer
Gizmo	\$19.99	Gadgets	GizmoWorks
Powergizmo	\$29.99	Gadgets	GizmoWorks
SingleTouch	\$149.99	Photography	Canon
MultiTouch	\$203.99	Household	Hitachi

```
SELECT *  
FROM Product  
WHERE category= 'Gadgets'
```



“selection”

PName	Price	Category	Manufacturer
Gizmo	\$19.99	Gadgets	GizmoWorks
Powergizmo	\$29.99	Gadgets	GizmoWorks

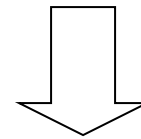


Simple SQL Query

Product

PName	Price	Category	Manufacturer
Gizmo	\$19.99	Gadgets	GizmoWorks
Powergizmo	\$29.99	Gadgets	GizmoWorks
SingleTouch	\$149.99	Photography	Canon
MultiTouch	\$203.99	Household	Hitachi

```
SELECT PName, Price, Manufacturer
FROM   Product
WHERE  Price > 100
```



“selection” and
“projection”

PName	Price	Manufacturer
SingleTouch	\$149.99	Canon
MultiTouch	\$203.99	Hitachi

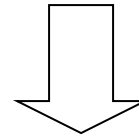
+ Notation



Input Schema

Product(PName, Price, Category, Manufacturer)

```
SELECT PName, Price, Manufacturer
FROM   Product
WHERE  Price > 100
```



Answer(PName, Price, Manufacturer)

Output Schema



Details



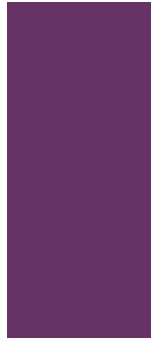
- Case insensitive:
 - Same: `SELECT` `Select` `select`
 - Same: `Product` `product`
 - Different: `'Seattle'` `'seattle'`
- Constants:
 - `'abc'` - yes
 - `"abc"` - no

+ The **LIKE** operator

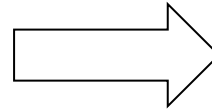
```
SELECT *  
FROM   Products  
WHERE  PName LIKE '%gizmo%'
```

- s **LIKE** p: pattern matching on strings
- p may contain two special symbols:
 - % = any sequence of characters
 - _ = any single character

+ Eliminating Duplicates



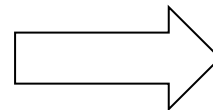
```
SELECT DISTINCT category  
FROM Product
```



Category
Gadgets
Photography
Household

Compare to:

```
SELECT category  
FROM Product
```



Category
Gadgets
Gadgets
Photography
Household

+ Ordering the Results



```
SELECT pname, price, manufacturer
FROM Product
WHERE category= 'gizmo' AND price > 50
ORDER BY price, pname
```

You can use more than one column in the ORDER BY clause.

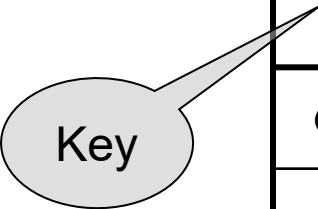
Make sure whatever column you are using to sort is in the column-list of the SELECT

Ties are broken by the second attribute on the ORDER BY list, etc.

Ordering is ascending, unless you specify the DESC keyword.
ORDER BY price ASC, pname DESC

+ Keys and Foreign Keys

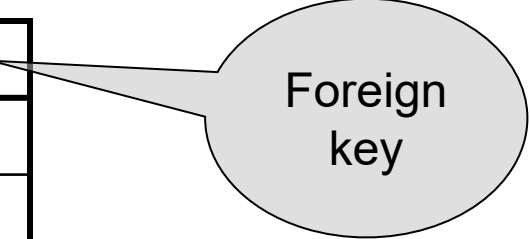
Company



<u>CName</u>	StockPrice	Country
GizmoWorks	25	USA
Canon	65	Japan
Hitachi	15	Japan

Product

<u>PName</u>	Price	Category	Manufacturer
Gizmo	\$19.99	Gadgets	GizmoWorks
Powergizmo	\$29.99	Gadgets	GizmoWorks
SingleTouch	\$149.99	Photography	Canon
MultiTouch	\$203.99	Household	Hitachi




Foreign
key

+ Joins

Product (pname, price, category, manufacturer)

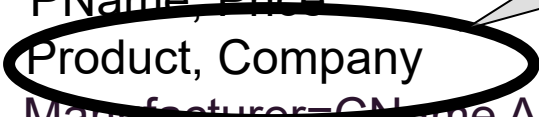
Company (cname, stockPrice, country)

Find all products under \$200 manufactured in Japan;
return their names and prices.



Join
between Product
and Company

```
SELECT PName, Price
FROM Product, Company
WHERE Manufacturer=CName AND Country= 'Japan'
AND Price <= 200
```



+ Joins

Product

PName	Price	Category	Manufacturer
Gizmo	\$19.99	Gadgets	GizmoWorks
Powergizmo	\$29.99	Gadgets	GizmoWorks
SingleTouch	\$149.99	Photography	Canon
MultiTouch	\$203.99	Household	Hitachi

Company

Cname	StockPrice	Country
GizmoWorks	25	USA
Canon	65	Japan
Hitachi	15	Japan

```
SELECT PName, Price
FROM Product, Company
WHERE Manufacturer=CName AND Country= 'Japan'
AND Price <= 200
```

PName	Price
SingleTouch	\$149.99

+ A Subtlety about Joins

Product (pname, price, category, manufacturer)

Company (cname, stockPrice, country)

Find all countries that manufacture some product in the 'Gadgets' category.

```
SELECT Country
FROM   Product, Company
WHERE  Manufacturer=CName AND Category= 'Gadgets'
```

Unexpected duplicates

+ A Subtlety about Joins

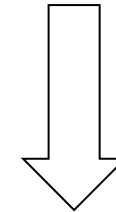
Product

<u>Name</u>	Price	Category	Manufacturer
Gizmo	\$19.99	Gadgets	GizmoWorks
Powergizmo	\$29.99	Gadgets	GizmoWorks
SingleTouch	\$149.99	Photography	Canon
MultiTouch	\$203.99	Household	Hitachi

Company

<u>Cname</u>	StockPrice	Country
GizmoWorks	25	USA
Canon	65	Japan
Hitachi	15	Japan

```
SELECT Country
FROM Product, Company
WHERE Manufacturer=CName AND Category='Gadgets'
```



Country
??
??

What is
the problem ?
What's the
solution ?



Tuple Variables



Person(pname, address, worksfor)
Company(cname, address)

```
SELECT DISTINCT pname, address
FROM   Person, Company
WHERE  worksfor = cname
```

Which
address ?



```
SELECT DISTINCT Person.pname, Company.address
FROM   Person, Company
WHERE  Person.worksfor = Company.cname
```



```
SELECT DISTINCT x.pname, y.address
FROM   Person AS x, Company AS y
WHERE  x.worksfor = y.cname
```



Subqueries Returning Relations

Company(name, city)

Product(pname, maker)

Purchase(id, product, buyer)

Return cities where one can find companies that manufacture products bought by Joe Blow

```
SELECT Company.city
FROM   Company
WHERE  Company.name IN
      (SELECT Product.maker
       FROM   Purchase , Product
       WHERE  Product.pname=Purchase.product
              AND Purchase .buyer = 'Joe Blow');
```



Subqueries Returning Relations



Is it equivalent to this ?

```
SELECT Company.city
FROM    Company, Product, Purchase
WHERE   Company.name= Product.maker
        AND Product.pname = Purchase.product
        AND Purchase.buyer = 'Joe Blow'
```

Beware of duplicates !

+ Removing Duplicates

```
SELECT DISTINCT Company.city
FROM   Company
WHERE  Company.name IN
      (SELECT Product.maker
       FROM   Purchase , Product
       WHERE  Product.pname=Purchase.product
              AND Purchase .buyer = 'Joe Blow');
```

```
SELECT DISTINCT Company.city
FROM   Company, Product, Purchase
WHERE  Company.name= Product.maker
      AND Product.pname = Purchase.product
      AND Purchase.buyer = 'Joe Blow'
```

Now
they are
equivalent



Subqueries Returning Relations

You can also use: $s > \text{ALL } R$
 $s > \text{ANY } R$
 $\text{EXISTS } R$

Product (pname, price, category, maker)

Find products that are more expensive than all those produced
By “Gizmo-Works”

```
SELECT name
FROM   Product
WHERE  price > ALL (SELECT price
                    FROM   Purchase
                    WHERE  maker= 'Gizmo-Works' )
```



Question: How are EXISTS and IN different?



- <http://dev.mysql.com/doc/refman/5.0/en/comparison-operators.html>
- http://dev.mysql.com/doc/refman/5.0/en/comparison-operators.html#function_in
- <http://dev.mysql.com/doc/refman/5.0/en/any-in-some-subqueries.html>
- <http://dev.mysql.com/doc/refman/5.1/en/exists-and-not-exists-subqueries.html>

Short answer: EXISTS returns TRUE if there is a row returned in the subquery, IN evaluates as TRUE if a value matches a value in a list (list of constants or subquery results)

+ Correlated Queries

Movie (title, year, director, length)

Find movies whose title appears more than once.

```
SELECT DISTINCT title
FROM   Movie AS x
WHERE  year <> ANY
      (SELECT year
       FROM   Movie
       WHERE  title = x.title);
```

correlation





Complex Correlated Query



Product (pname, price, category, maker, year)

- Find products (and their manufacturers) that are more expensive than all products made by the same manufacturer before 1972

```
SELECT DISTINCT pname, maker
FROM   Product AS x
WHERE  price > ALL (SELECT price
                   FROM   Product AS y
                   WHERE  x.maker = y.maker AND y.year < 1972);
```



Aggregation

```
SELECT avg(price)
FROM Product
WHERE maker="Toyota"
```

```
SELECT count(*)
FROM Product
WHERE year > 1995
```

SQL supports several aggregation operations:

sum, count, min, max, avg

Except count, all aggregations apply to a single attribute

+ Aggregation: Count

COUNT applies to duplicates, unless otherwise stated:

```
SELECT Count(category)
FROM   Product
WHERE  year > 1995
```

same as Count(*)

We probably want:

```
SELECT Count(DISTINCT category)
FROM   Product
WHERE  year > 1995
```



More Examples



Purchase(product, date, price, quantity)

```
SELECT Sum(price * quantity)
FROM Purchase
```

```
SELECT Sum(price * quantity)
FROM Purchase
WHERE product = 'bagel'
```

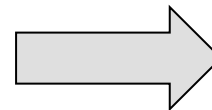
What do
they mean ?

+ Simple Aggregations

Purchase

Product	Date	Price	Quantity
Bagel	10/21	1	20
Banana	10/3	0.5	10
Banana	10/10	1	10
Bagel	10/25	1.50	20

```
SELECT Sum(price * quantity)
FROM Purchase
WHERE product = 'bagel'
```



50 (= 20+30)



Grouping and Aggregation



Purchase(product, date, price, quantity)

Find total sales after 10/1/2005 per product.

```
SELECT    product, Sum(price*quantity) AS TotalSales
FROM      Purchase
WHERE     date > '10/1/2005'
GROUP BY  product
```

Let's see what this means...



Grouping and Aggregation



1. Compute the **FROM** and **WHERE** clauses.
2. Group by the attributes in the **GROUPBY**
3. Compute the **SELECT** clause: grouped attributes and aggregates.



1&2. FROM-WHERE-GROUPBY

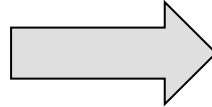


Product	Date	Price	Quantity
Bagel	10/21	1	20
Bagel	10/25	1.50	20
Banana	10/3	0.5	10
Banana	10/10	1	10



3. SELECT

Product	Date	Price	Quantity
Bagel	10/21	1	20
Bagel	10/25	1.50	20
Banana	10/3	0.5	10
Banana	10/10	1	10



Product	TotalSales
Bagel	50
Banana	15

```
SELECT    product, Sum(price*quantity) AS TotalSales
FROM      Purchase
WHERE     date > '10/1/2005'
GROUP BY product
```

GROUP BY v.s. Nested Quereis

```
SELECT    product, Sum(price*quantity) AS TotalSales
FROM      Purchase
WHERE     date > '10/1/2005'
GROUP BY  product
```

```
SELECT DISTINCT x.product, (SELECT Sum(y.price*y.quantity)
                             FROM    Purchase y
                             WHERE   x.product = y.product
                             AND y.date > '10/1/2005' )
                             AS TotalSales
FROM      Purchase x
WHERE     x.date > '10/1/2005'
```



Another Example

What does
it mean ?

```
SELECT    product,  
          sum(price * quantity) AS SumSales  
          max(quantity) AS MaxQuantity  
FROM      Purchase  
GROUP BY product
```



HAVING Clause



Same query, except that we consider only products that had at least 100 buyers.

```
SELECT    product, Sum(price * quantity)
FROM      Purchase
WHERE     date > '10/1/2005'
GROUP BY  product
HAVING    Sum(quantity) > 30
```

HAVING clause contains conditions on aggregates.



General form of Grouping and Aggregation



```
SELECT  S
FROM    R1,...,Rn
WHERE   C1
GROUP BY a1,...,ak
HAVING  C2
```



S = may contain attributes a_1, \dots, a_k and/or any aggregates but NO OTHER ATTRIBUTES

C1 = is any condition on the attributes in R_1, \dots, R_n

C2 = is any condition on aggregate expressions



General form of Grouping and Aggregation

```
SELECT  S
FROM    R1,...,Rn
WHERE   C1
GROUP BY a1,...,ak
HAVING  C2
```

Evaluation steps:

1. Evaluate FROM-WHERE, apply condition C1
2. Group by the attributes a_1, \dots, a_k
3. Apply condition C2 to each group (may have aggregates)
4. Compute aggregates in S and return the result

+ Advanced SQLizing



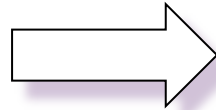
1. Getting around INTERSECT and EXCEPT
2. Quantifiers
3. Aggregation v.s. subqueries



INTERSECT and EXCEPT: not in SQL Server

1. INTERSECT and EXCEPT:

```
(SELECT R.A, R.B  
FROM R)  
INTERSECT  
(SELECT S.A, S.B  
FROM S)
```



```
SELECT R.A, R.B  
FROM R  
WHERE  
EXISTS(SELECT *  
FROM S  
WHERE R.A=S.A and R.B=S.B)
```

If R, S have no
duplicates, then can
write without
subqueries
(HOW ?)

```
(SELECT R.A, R.B  
FROM R)  
EXCEPT  
(SELECT S.A, S.B  
FROM S)
```



```
SELECT R.A, R.B  
FROM R  
WHERE  
NOT EXISTS(SELECT *  
FROM S  
WHERE R.A=S.A and R.B=S.B)
```



2. Quantifiers



Product (pname, price, company)
Company(cname, city)

Find all companies that make some products with price < 100

```
SELECT DISTINCT Company.cname  
FROM    Company, Product  
WHERE   Company.cname = Product.company and Product.price < 100
```

Existential: easy ! 😊



2. Quantifiers



Product (pname, price, company)
Company(cname, city)

Find all companies that make only products with price < 100

same as:

Find all companies s.t. all of their products have price < 100

Universal: hard ! 😞



2. Quantifiers



1. Find *the other* companies: i.e. s.t. some product ≥ 100

```
SELECT DISTINCT Company.cname
FROM   Company
WHERE  Company.cname IN (SELECT Product.company
                        FROM Product
                        WHERE Produc.price  $\geq$  100)
```

2. Find all companies s.t. all their products have price < 100

```
SELECT DISTINCT Company.cname
FROM   Company
WHERE  Company.cname NOT IN (SELECT Product.company
                            FROM Product
                            WHERE Produc.price  $\geq$  100)
```



3. Group-by v.s. Nested Query

Author(login,name)

Wrote(login,url)

- Find authors who wrote ≥ 10 documents:
- Attempt 1: with nested queries

```
SELECT DISTINCT Author.name
FROM      Author
WHERE     count(SELECT Wrote.url
                  FROM Wrote
                  WHERE Author.login=Wrote.login)
          > 10
```

This is
SQL by
a novice




3. Group-by v.s. Nested Query



- Find all authors who wrote at least 10 documents:
- Attempt 2: SQL style (with GROUP BY)

```
SELECT    Author.name
FROM      Author, Wrote
WHERE     Author.login=Wrote.login
GROUP BY  Author.name
HAVING    count(wrote.url) > 10
```



This is
SQL by
an expert

No need for **DISTINCT**: automatically from **GROUP BY**



3. Group-by v.s. Nested Query

Author(login,name)

Wrote(login,url)

Mentions(url,word)

Find authors with vocabulary ≥ 10000 words:

```
SELECT    Author.name
FROM      Author, Wrote, Mentions
WHERE     Author.login=Wrote.login AND Wrote.url=Mentions.url
GROUP BY  Author.name
HAVING    count(distinct Mentions.word) > 10000
```



Two Examples



Store(sid, sname)

Product(pid, pname, price, sid)

Find all stores that sell *only* products with price > 100

same as:

Find all stores s.t. all their products have price > 100)


```
SELECT Store.name
FROM   Store, Product
WHERE  Store.sid = Product.sid
GROUP BY Store.sid, Store.name
HAVING 100 < min(Product.price)
```

Why both ?

Almost equivalent...

```
SELECT Store.name
FROM   Store
WHERE
    100 < ALL (SELECT Product.price
                FROM product
                WHERE Store.sid = Product.sid)
```

```
SELECT Store.name
FROM   Store
WHERE  Store.sid NOT IN
      (SELECT Product.sid
       FROM Product
       WHERE Product.price <= 100)
```



Two Examples



Store(sid, sname)

Product(pid, pname, price, sid)

For each store,
find its most expensive product



Two Examples

This is easy but doesn't do what we want:

```
SELECT Store.sname, max(Product.price)
FROM   Store, Product
WHERE  Store.sid = Product.sid
GROUP BY Store.sid, Store.sname
```

Better:

```
SELECT Store.sname, x.pname
FROM   Store, Product x
WHERE  Store.sid = x.sid and
      x.price >=
      ALL (SELECT y.price
            FROM Product y
            WHERE Store.sid = y.sid)
```

But may
return
multiple
product names
per store



Two Examples

Finally, choose some pid arbitrarily, if there are many with highest price:

```
SELECT Store.sname, max(x.pname)
FROM   Store, Product x
WHERE  Store.sid = x.sid and
       x.price >=
       ALL (SELECT y.price
            FROM Product y
            WHERE Store.sid = y.sid)
GROUP BY Store.sname
```



Announcement



- “A note taker is required to assist a student in this class. There is an honorarium of \$75/course/term, with some conditions. If you are interested, please go to the Advising and Access Services Centre (AASC), 6227 University Avenue, to obtain information and to fill out a Confidentiality form.”
- I wanted to confirm for you that the notetaker would be paid in full and will mostly likely be asked to provide retroactive notes for the term.



Subqueries Returning Relations

You can also use: $s > \text{ALL } R$
 $s > \text{ANY } R$
 $\text{EXISTS } R$

Product (pname, price, category, maker)

Find products that are more expensive than all those produced
By “Gizmo-Works”

```
SELECT name
FROM   Product
WHERE  price > ALL (SELECT price
                    FROM   Purchase
                    WHERE  maker= 'Gizmo-Works' )
```



Question: How are EXISTS and IN different?



- <http://dev.mysql.com/doc/refman/5.0/en/comparison-operators.html>
- http://dev.mysql.com/doc/refman/5.0/en/comparison-operators.html#function_in
- <http://dev.mysql.com/doc/refman/5.0/en/any-in-some-subqueries.html>
- <http://dev.mysql.com/doc/refman/5.1/en/exists-and-not-exists-subqueries.html>

Short answer: EXISTS returns TRUE if there is a row returned in the subquery, IN evaluates as TRUE if a value matches a value in a list (list of constants or subquery results)



NULLS in SQL



- Whenever we do not have a value, we can put a NULL
- Can mean many things:
 - Value does not exist
 - Value exists but is unknown
 - Value not applicable
 - Etc.
- The schema specifies for each attribute if it can be null (*nullable* attribute) or not
- How does SQL cope with tables that have NULLs ?



Null Values



- If $x = \text{NULL}$ then $4 \cdot (3 - x) / 7$ is still **NULL**
- If $x = \text{NULL}$ then $x = \text{"Joe"}$ is **UNKNOWN**
- In SQL there are three boolean values:

FALSE	=	0
UNKNOWN	=	0.5
TRUE	=	1



Null Values



- $C1 \text{ AND } C2 = \min(C1, C2)$
- $C1 \text{ OR } C2 = \max(C1, C2)$
- $\text{NOT } C1 = 1 - C1$

```
SELECT *  
FROM Person  
WHERE (age < 25) AND  
      (height > 6 OR weight > 190)
```

E.g.
age=20
height=NULL
weight=200

Rule in SQL: include only tuples that yield TRUE



Null Values



Unexpected behavior:

```
SELECT *  
FROM    Person  
WHERE   age < 25 OR age >= 25
```

Some Persons are not included !



Null Values



Can test for NULL explicitly:

- x IS NULL
- x IS NOT NULL

```
SELECT *  
FROM   Person  
WHERE  age < 25 OR age >= 25 OR  
       age IS NULL
```

Now it includes all Persons



Outerjoins



Explicit joins in SQL = “inner joins” :

Product(name, category)

Purchase(prodName, store)

```
SELECT Product.name, Purchase.store
FROM   Product JOIN Purchase ON
        Product.name = Purchase.prodName
```

Same as:

```
SELECT Product.name, Purchase.store
FROM   Product, Purchase
WHERE  Product.name = Purchase.prodName
```

But Products that never sold will be lost !



Outer Joins



- Left outer join:

- Include the left tuple even if there's no match

- Right outer join:

- Include the right tuple even if there's no match

- Full outer join:

- Include the both left and right tuples even if there's no match



Outerjoins



Left outer joins in SQL:

Product(name, category)

Purchase(prodName, store)

```
SELECT Product.name, Purchase.store
FROM   Product LEFT OUTER JOIN Purchase ON
        Product.name = Purchase.prodName
```



```
SELECT Product.name, Purchase.store
FROM   Product LEFT OUTER JOIN Purchase ON
        Product.name = Purchase.prodName
```

Product

Name	Category
Gizmo	gadget
Camera	Photo
OneClick	Photo

Purchase

ProdName	Store
Gizmo	Wiz
Camera	Ritz
Camera	Wiz

Name	Store
Gizmo	Wiz
Camera	Ritz
Camera	Wiz
OneClick	NULL



Application



Compute, for each product, the total number of sales in
'September'

Product(name, category)

Purchase(prodName, month, store)

```
SELECT Product.name, count(*)  
FROM   Product, Purchase  
WHERE  Product.name = Purchase.prodName  
       and Purchase.month = 'September'  
GROUP BY Product.name
```

What's wrong ?



Application



Compute, for each product, the total number of sales in
'September'

Product(name, category)

Purchase(prodName, month, store)

```
SELECT Product.name, count(*)  
FROM   Product LEFT OUTER JOIN Purchase ON  
        Product.name = Purchase.prodName  
        and Purchase.month = 'September'  
GROUP BY Product.name
```

Now we also get the products who sold in 0 quantity



Modifying the Database



Three kinds of modifications

- Insertions
- Deletions
- Updates

Sometimes they are all called “updates”

+ Insertions

General form:

```
INSERT INTO R(A1,..., An) VALUES (v1,..., vn)
```

Example: Insert a new purchase to the database:

```
INSERT INTO Purchase(buyer, seller, product, store)
VALUES ( 'Joe' , 'Fred' , 'wakeup-clock-espresso-machine' ,
        'The Sharper Image' )
```

Missing attribute → NULL.

May drop attribute names if give them in order.



Insertions



```
INSERT INTO PRODUCT(name)

  SELECT DISTINCT Purchase.product
  FROM   Purchase
  WHERE  Purchase.date > "10/26/01"
```

The query replaces the VALUES keyword.
Here we insert *many* tuples into PRODUCT

+ Insertion: an Example

Product(name, listPrice, category)
Purchase(prodName, buyerName, price)

prodName is foreign key in Product.name

Suppose database got corrupted and we need to fix it:

Product

name	listPrice	category
gizmo	100	gadgets

Purchase

prodName	buyerName	price
camera	John	200
gizmo	Smith	80
camera	Smith	225

Task: insert in Product all prodNames from Purchase

+ Insertion: an Example

```
INSERT INTO Product(name)

SELECT DISTINCT prodName
FROM Purchase
WHERE prodName NOT IN (SELECT name FROM Product)
```

name	listPrice	category
gizmo	100	Gadgets
camera	-	-

+ Insertion: an Example

```
INSERT INTO Product(name, listPrice)
```

```
SELECT DISTINCT prodName, price
```

```
FROM Purchase
```

```
WHERE prodName NOT IN (SELECT name FROM Product)
```

name	listPrice	category
gizmo	100	Gadgets
camera	200	-
camera ??	225 ??	-

← Depends on the implementation



Deletions



Example:

```
DELETE FROM PURCHASE  
  
WHERE seller = 'Joe' AND  
       product = 'Brooklyn Bridge'
```

Factoid about SQL: there is no way to delete only a single occurrence of a tuple that appears twice in a relation.

+ Updates



Example:

```
UPDATE PRODUCT
SET price = price/2
WHERE Product.name IN
      (SELECT product
       FROM Purchase
       WHERE Date = 'Oct, 25, 1999');
```